

# 5. Feature Tracking

Homography estimation (previous lecture), Optical Flow, Feature Tracking, Motion Segmentation

# Estimating transformations from Image features

- If point correspondences  $(x,y) \leftrightarrow (x',y')$  are known


We can find a 3X3 transformation,  $\mathbf{A}$  such that  $\mathbf{x}' = \mathbf{A}\mathbf{x}$

$$\mathbf{A} = [a_1, a_2, a_3; a_4, a_5, a_6; a_7, a_8, a_9]$$


Remember from first lecture on transformations, a 2D projective transformation (most generic transformation) has 8 dof

Assume  $a_9 = 1$ , therefore

$$x' = \frac{a_1x + a_2y + a_3}{a_7x + a_8y + 1}$$
$$y' = \frac{a_4x + a_5y + a_6}{a_7x + a_8y + 1}$$




# Feature Matching Performance



The distance threshold affects performance

- True positives = # of detected matches that are correct
- Suppose we want to maximize these—how to choose threshold?
- False positives = # of detected matches that are incorrect
- Suppose we want to minimize these—how to choose threshold?

$$\frac{\text{\# true positives}}{\text{\# correctly matched features (positives)}}$$
  
"recall"



# Estimating transformations from Image features

Assume  $a_9 = 1$ , therefore

$$x' = \frac{a_1x + a_2y + a_3}{a_7x + a_8y + 1}$$

$$y' = \frac{a_4x + a_5y + a_6}{a_7x + a_8y + 1}$$


$$a_7x'x + a_8x'y + x' = a_1x + a_2y + a_3$$

$$a_7y'x + a_8y'y + y' = a_4x + a_5y + a_6$$

$$x' = a_1x + a_2y + a_3 - a_7x'x - a_8x'y$$

$$y' = a_4x + a_5y + a_6 - a_7y'x - a_8y'y$$

Segregate the equations into matrices/vectors of knowns and unknowns



# Estimating transformations from Image features

$$a_7x'x + a_8x'y + x' = a_1x + a_2y + a_3$$

$$a_7y'x + a_8y'y + y' = a_7x + a_8y + a_6$$

$$x' = a_1x + a_2y + a_3 - a_7x'x - a_8x'y$$


$$y' = a_7x + a_8y + a_6 - a_7y'x - a_8y'y$$

Two rows for each point  $i$

$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & 0 & -x_ix'_i & -y_ix'_i \\ 0 & 0 & 0 & x_i & y_i & 1 & -x_iy'_i & -y_iy'_i \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \\ a_5 \\ a_6 \\ a_7 \\ a_8 \end{bmatrix} = \begin{bmatrix} x'_i \\ y'_i \end{bmatrix}$$

$$\mathbf{T}_{8 \times 8} \mathbf{a}_{8 \times 1} = \mathbf{B}_{8 \times 1}$$

This looks like Linear Least squares formulation. Once the solution is obtained re-arrange the solution to form  $\mathbf{A}$





# Homography using Direct Linear Transformation (DLT)

Don't assume  $a_9 = 1$ , add it to the system

In order to obtain a non-trivial solution,  
we obtain  $\mathbf{a}$  :

$$\min \|\mathbf{T}'\mathbf{a}\| \text{ such that } \|\mathbf{a}\| = 1$$

Two rows for each point  $i$

$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & 0 & -x_i x'_i & -y_i x'_i \\ 0 & 0 & 0 & x_i & y_i & 1 & -x_i y'_i & -y_i y'_i \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \\ a_5 \\ a_6 \\ a_7 \\ a_8 \end{bmatrix} = \begin{bmatrix} \vdots \\ x'_i \\ y'_i \\ \vdots \end{bmatrix}$$

$\mathbf{a}_i$

$$\mathbf{T}'_{8 \times 9} \mathbf{a}_{9 \times 1} = \mathbf{0}_{9 \times 1}$$

$$\mathbf{T}' = [\mathbf{T}, -\mathbf{B}]$$

# Homography using Direct Linear Transformation (DLT)

Don't assume  $a_9 = 1$ , add it to the system

In order to obtain a non-trivial solution,  
we obtain  $\mathbf{a}$  :

$$\min \|\mathbf{T}'\mathbf{a}\| \text{ such that } \|\mathbf{a}\| = 1$$

Solution: SVD of  $\mathbf{T}'$ . The eigenvector corresponding to smallest eigenvalue is the solution.

In order to get a stable solution, Normalise the data before DLT.

- a) Translate for zero mean
- b) Scale so that average distance to origin is  $\sim \sqrt{2}$

[https://www.cs.cmu.edu/~16385/s17/Slides/10.2\\_2D\\_Alignment\\_DLT.pdf](https://www.cs.cmu.edu/~16385/s17/Slides/10.2_2D_Alignment_DLT.pdf)

# Homography using Direct Linear Transformation (DLT)

Don't assume  $a_9 = 1$ , add it to the system

In order to obtain a non-trivial solution,  
we obtain  $\mathbf{a}$  :

$$\min \|\mathbf{T}'\mathbf{a}\| \text{ such that } \|\mathbf{a}\| = 1$$

Solution: SVD of  $\mathbf{T}'$ . The eigenvector corresponding to smallest eigenvalue is the solution.

In order to get a stable solution, Normalise the data before DLT.


- a) Translate for zero mean
- b) Scale so that average distance to origin is  $\sim \sqrt{2}$

[https://www.cs.cmu.edu/~16385/s17/Slides/10.2\\_2D\\_Alignment\\_DLT.pdf](https://www.cs.cmu.edu/~16385/s17/Slides/10.2_2D_Alignment_DLT.pdf)




# Panoramic Image Stitching



Given a set of images



Obtain:





# In Practice, we map planes




For red image: pixels are already on the plane  
For green image: map to first image plane

# Planar mapping



For red image: pixels are already on the plane  
For green image: map to first image plane

# Solution: Use cylindrical mapping instead



For red image: compute  $h$ ,  $\theta$  on cylindrical surface from  $(u, v)$

For green image: map to first image plane, then map to cylindrical surface



# Solution: Use cylindrical mapping instead

Calculate angle and height:

$$\theta = (x - x_c) / f$$

$$h = (y - y_c) / f$$

Find unit cylindrical coords:

$$X' = \sin(\theta)$$


$$Y' = h$$

$$Z' = \cos(\theta)$$

Project to image plane:

$$x' = f X' / Z' + x_c$$

$$y' = f Y' / Z' + y_c$$



# RANSAC

Look for ‘only’ inliers.

1. Randomly select a seed group to estimate transformation
2. Compute transformation
3. Compute inliers
4. If inliers are large, recompute transformation

Keep transformation with most inliers

For octave, you can use <https://github.com/RANSAC/RANSAC-Toolbox>

# P1: Image warping

1. Warp simpsons.jpeg to bus.jpeg such that simpsons appear on the bus advertisement

You can manually select the image features for matching



# P2: Image stitching

1. Build a panorama using keble images.
2. Use at least 3 methods to find features.
3. Estimate transformations with and without RANSAC
4. Manually match images and estimate transformation
5. Is any result from 3 is close to 4 ?
6. Repeat the experiment on various images of “La Place des Jacobins”.

Bonus: Can you automatically align images?

# Feature Tracking & Optical Flow


Feature tracking: Extract visual features (corners, textured areas) and “track” them over multiple frames

Optical Flow: Recover image motion at each pixel from spatio-temporal image brightness variations

Optical flow = apparent motion of brightness patterns

# Optical Flow

Why bother with motion?



G. Johansson, "Visual Perception of Biological Motion and a Model For Its Analysis", *Perception and Psychophysics* 14, 201-211, 1973.

# Feature Tracking & Optical Flow

Challenges:

1. Figure out which features can be tracked
2. Efficiently track across frames
3. Some points may change appearance over time (e.g., due to rotation, moving into shadows, etc.)
4. Drift: small errors can accumulate as appearance model is updated
5. Points may appear or disappear: need to be able to add/delete tracked points

# Feature Tracking & Optical Flow

Feature tracking: Extract visual features (corners, textured areas) and “track” them over multiple frames

Problems: Sparse, inexact feature alignment, low accuracy

Optical Flow: Recover image motion at each pixel from spatio-temporal image brightness variations


Advantage: Scale/rotation invariant, lighting invariant, can handle large movements (not really, not accurately for sure)

Registration using Lucas-Kanade Method


[An iterative image registration technique with an application to stereo vision](#)

# Feature Tracking & Optical Flow


Feature-tracking



Optical flow from the side window of a car




Optical flow for a pilot landing a plane




Applications:

3D reconstruction, Object segmentation, Motion Segmentation, Tracking, Super-resolution, Activity recognition

# Optical Flow



$I(x,y,t)$



$I(x,y,t+1)$


Given two subsequent frames, estimate the point translation

Key assumptions of Lucas-Kanade Tracker

- **Brightness constancy:** projection of the same point looks the same in every frame
- **Small motion:** points do not move very far
- **Spatial coherence:** points move like their neighbors



# Optical Flow: Brightness Constancy



Brightness Constancy Equation:  $I(x, y, t) = I(x + u, y + v, t + 1)$

Take Taylor expansion of  $I(x + u, y + v, t + 1)$  at  $(x, y, t)$  to linearize the right side:

$$I(x + u, y + v, t + 1) \approx I(x, y, t) + \overset{\text{Image derivative along x}}{I_x} \cdot u + I_y \cdot v + \overset{\text{Difference over frames}}{I_t}$$
$$I(x + u, y + v, t + 1) - I(x, y, t) = I_x \cdot u + I_y \cdot v + I_t$$

So:  $I_x \cdot u + I_y \cdot v + I_t \approx 0 \rightarrow \nabla I \cdot \begin{bmatrix} u & v \end{bmatrix}^T + I_t = 0$

# Optical Flow: Brightness Constancy


$$\nabla I \cdot [u \ v]^T + I_t = 0$$

1 equation, 2 unknowns (u,v), can't solve


The component of the motion perpendicular to the gradient (i.e., parallel to the edge) cannot be measured

If  $(u, v)$  satisfies the equation,  
so does  $(u+u', v+v')$  if


$$\nabla I \cdot [u' \ v']^T = 0$$



# Optical Flow: Aperture Problem




Actual motion




Perceived motion

# Optical Flow: Aperture Problem



Actual motion



Perceived motion

Solution:  
Look at neighbors

# Optical Flow: Barber Pole Illusion



[http://en.wikipedia.org/wiki/Barberpole\\_illusion](http://en.wikipedia.org/wiki/Barberpole_illusion)

# Optical Flow: Barber Pole Illusion



[http://en.wikipedia.org/wiki/Barberpole\\_illusion](http://en.wikipedia.org/wiki/Barberpole_illusion)

# Optical Flow: Solving the issue

Get more equations for a pixel.

Spatial coherence constraint: Assume the pixel's neighbors have the same (u,v)

If we use a 5x5 window, that gives us 25 equations per pixel

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v]$$
$$\begin{bmatrix} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{bmatrix} \quad \Rightarrow \quad \begin{matrix} A & d = b \\ 25 \times 2 & 2 \times 1 & 25 \times 1 \\ (A^T A) & d = A^T b \end{matrix}$$



# Optical Flow: Solving the issue

When is this solvable? I.e., what are good points to track? Remember: Harris Corners

- $\mathbf{A}^T\mathbf{A}$  should be invertible
- $\mathbf{A}^T\mathbf{A}$  should not be too small due to noise
  - eigenvalues  $\lambda_1$  and  $\lambda_2$  of  $\mathbf{A}^T\mathbf{A}$  should not be too small
- $\mathbf{A}^T\mathbf{A}$  should be well-conditioned
  - $\lambda_1 / \lambda_2$  should not be too large ( $\lambda_1 =$  larger eigenvalue)

Low-texture region




- gradients are small magnitude
- small eigenvalues

Edge



- gradients are large and small
- large and small eigenvalue

High-texture region




- gradients are different, large magnitudes
- large eigenvalues

# Optical Flow- Lucas-Kanade

Assumptions:

1. Brightness consistency
2. Spatial coherence
3. Small motion (not really!)

# Dealing with larger movements: coarse-to-fine registration



# Optical Flow: State of the Art



Start with something similar to Lucas-Kanade

+ gradient constancy

+ energy minimization with smoothing term

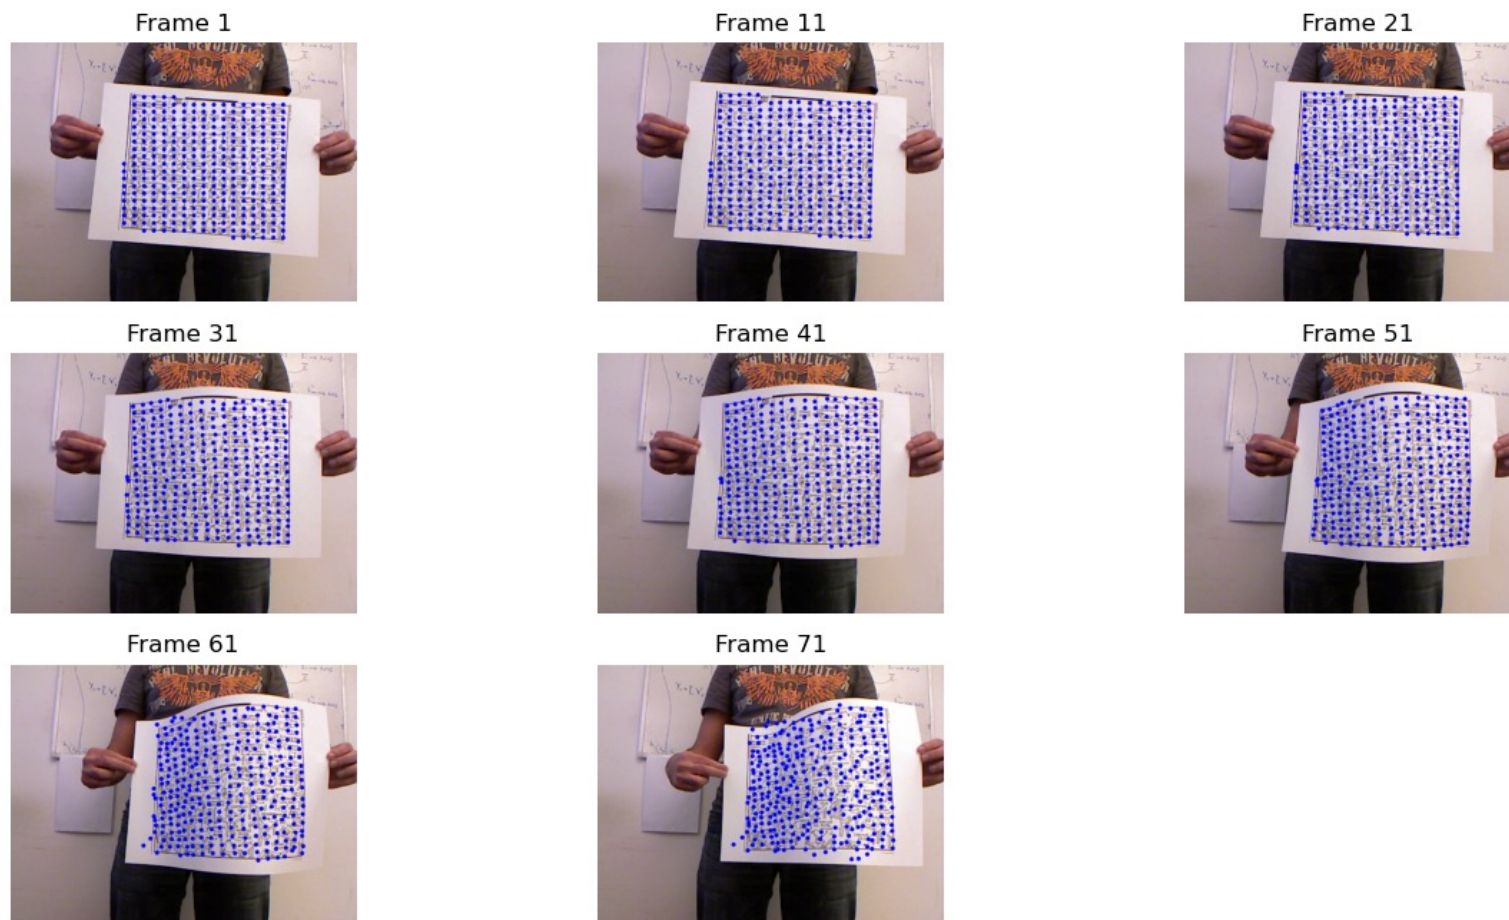
+ region matching

+ keypoint matching (long-range)



Region-based +Pixel-based +Keypoint-based

# Optical Flow: State of the Art



# Feature Tracking: Shi-Tomasi Feature Tracker

Find good features using eigenvalues of second-moment matrix (e.g., Harris detector or threshold on the smallest eigenvalue)

Key idea: “good” features to track are the ones whose motion can be estimated reliably

Track from frame to frame with Lucas-Kanade

This amounts to assuming a translation model for frame-to-frame feature movement

Check consistency of tracks by *affine* registration to the first observed instance of the feature

Affine model is more accurate for larger displacements

Comparing to the first frame helps to minimize drift

# Feature Tracking: Shi-Tomasi Feature Tracker



Figure 1: Three frame details from Woody Allen's *Manhattan*. The details are from the 1st, 11th, and 21st frames of a subsequence from the movie.



Figure 2: The traffic sign windows from frames 1,6,11,16,21 as tracked (top), and warped by the computed deformation matrices (bottom).



# Feature Tracking: Shi-Tomasi Feature Tracker


1. Find a good point to track (harris corner)
2. Use intensity second moment matrix and difference across frames to find displacement
3. Iterate and use coarse-to-fine search to deal with larger movements
4. When creating long tracks, check appearance of registered patch against appearance of initial patch to find points that have drifted

## Implementation issues:

- Window size
  - Small window more sensitive to noise and may miss larger motions (without pyramid)
  - Large window more likely to cross an occlusion boundary (and it's slower)
  - 15x15 to 31x31 seems typical
- Weighting the window
  - Common to apply weights so that center matters more (e.g., with Gaussian)

# Motion Segmentation

- Create layers (with coherent affine motion) and track them




J. Wang and E. Adelson. [Layered Representation for Motion Analysis](#). CVPR 1993.

# Affine Motion

$$u(x, y) = a_1 + a_2x + a_3y$$

$$v(x, y) = a_4 + a_5x + a_6y$$



- Substituting into the brightness constancy equation:

$$I_x(a_1 + a_2x + a_3y) + I_y(a_4 + a_5x + a_6y) + I_t \approx 0$$


- Each pixel provides 1 linear constraint in 6 unknowns
- Least squares minimization:

$$Err(\underline{a}) = \sum \left[ I_x(a_1 + a_2x + a_3y) + I_y(a_4 + a_5x + a_6y) + I_t \right]^2$$

Source: Silvio Savarese

# How to estimate layers?

1. Obtain a set of initial affine motion hypotheses
  - Divide the image into blocks and estimate affine motion parameters in each block by least squares
  - Eliminate hypotheses with high residual error
2. Map into motion parameter space
3. Perform k-means clustering on affine motion parameters
  - Merge clusters that are close and retain the largest clusters to obtain a smaller set of hypotheses to describe all the motions in the scene



Source: Silvio Savarese

# How to estimate layers?


## 1. Obtain a set of initial affine motion hypotheses

- Divide the image into blocks and estimate affine motion parameters in each block by least squares
- Eliminate hypotheses with high residual error

## 2. Map into motion parameter space

## 3. Perform k-means clustering on affine motion parameters

- Merge clusters that are close and retain the largest clusters to obtain a smaller set of hypotheses to describe all the motions in the scene





# How to estimate layers?

## 1. Obtain a set of initial affine motion hypotheses


- Divide the image into blocks and estimate affine motion parameters in each block by least squares
- Eliminate hypotheses with high residual error
- Map into motion parameter space
- Perform k-means clustering on affine motion parameters
- Merge clusters that are close and retain the largest clusters to obtain a smaller set of hypotheses to describe all the motions in the scene

## 2. Iterate until convergence:

- Assign each pixel to best hypothesis
- Pixels with high residual error remain unassigned
- Perform region filtering to enforce spatial constraints
- Re-estimate affine motions in each region

Source: Silvio Savarese

# How to estimate layers?



J. Wang and E. Adelson. [Layered Representation for Motion Analysis](#). CVPR 1993.

Source: Silvio Savarese

# Course Project: Rubik's cube.

Use a webcam to detect all small squares, obtain cube orientation and reconstruct the cube.

Goals:

1. Detect a cube in the video stream through webcam. Be careful of multiple configurations. (30%)
2. Detect all small squares. The application should be able to count number of how many squares of each color are visible and how many faces are visible. (30%)
3. Detect the cube orientation. Display the information. (15%)
4. Generate 3D model (15%)

To submit: User manual, (OPENCV/MATLAB) code + GUI, well-documented code

Online Demo in first week of January